

Towards Ethical Intelligence: Navigating Fairness and Bias in AI

Roberta Calegari

Alma Mater Studiorum–Università di Bologna, Italy

BIAS 2023 — 3rd Workshop on Bias and Fairness in AI
Workshop at ECML PKDD 2023
Turin, Italy

22 September 2023



Funded by
the European Union

www.aequitas-project.eu
info@aequitas-project.eu

Next in Line...

- 1 Why fairness?
- 2 Goal of the talk
- 3 Fairness and bias in AI
- 4 Fairness Awareness
- 5 Enforcing Fairness
- 6 Our advancements in the field
- 7 Conclusions, Challenges and Opportunities



Why fairness?

- society is facing a dramatic increase in *pervasive inequality* and *intersectional discrimination* due to the widespread use of AI

[Leavy et al., 2021, Leavy et al., 2020]

- ML is contributing to creating a society where some groups or individuals are disadvantaged
- <https://www.propublica.org/article/machine-bias-risk-assessments-in-criminal-sentencing>
- <https://www.technologyreview.com/s/610634/microsofts-neo-nazi-sexbot-was-a-great-lesson-for-makers-of-ai-assistants/>



EG-TAI: TAI Requirements & AI Act

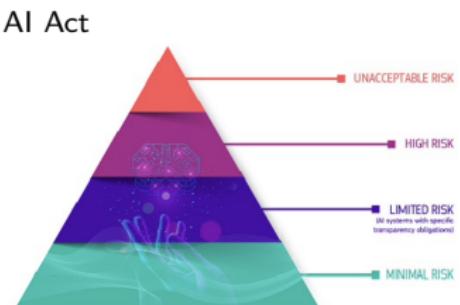
Main pillars

- lawfulness
- ethics
- robustness



Seven specific requirements – dimensions to be audited – of an AI system:

- ① human agency and oversight
- ② technical robustness and safety
- ③ privacy and data governance
- ④ transparency (traceability, explainability)
- ⑤ diversity, non-discrimination and *fairness*
- ⑥ societal and environmental well-being
- ⑦ accountability



Next in Line...

- 1 Why fairness?
- 2 Goal of the talk
- 3 Fairness and bias in AI
- 4 Fairness Awareness
- 5 Enforcing Fairness
- 6 Our advancements in the field
- 7 Conclusions, Challenges and Opportunities



Outline

- Fairness: state of the art (awareness/enforcement)
- Our advancements
- Challenges and opportunities



Next in Line...

- 1 Why fairness?
- 2 Goal of the talk
- 3 Fairness and bias in AI
- 4 Fairness Awareness
- 5 Enforcing Fairness
- 6 Our advancements in the field
- 7 Conclusions, Challenges and Opportunities



What is fairness? I

Article 21 of the EU Charter of Fundamental Rights

any discrimination based on any ground such as sex, race, colour, ethnic or social origin, genetic features, language, religion or belief, political or any other opinion, membership of a national minority, property, birth, disability, age or sexual orientation shall be prohibited.

They describe two different discrimination scenarios:

- ① direct discrimination (disparate *treatment*)
- ② indirect discrimination (disparate *impact*): when a seemingly “neutral provision, criterion or practice” disproportionately disadvantages members of a given sensitive group compared to others



What is bias? I

Bias and fairness in AI: two sides of the same coin

While there is no universally agreed upon definition for fairness, we can broadly define fairness as

the absence of prejudice or preference for an individual or group based on their characteristics, i.e., absence of bias

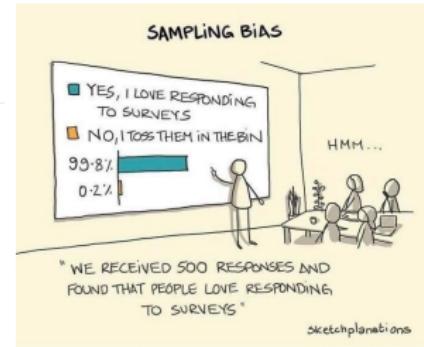
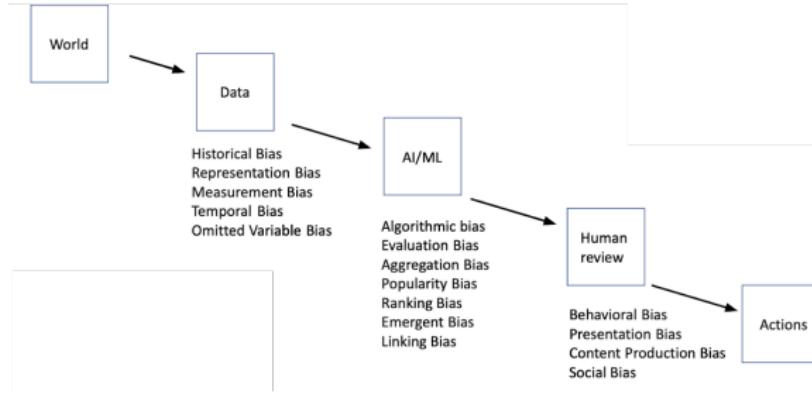
Bias in AI

Phenomenon that occurs when an AI system produces results that are systematically prejudiced

- many shapes and forms of bias
- can be introduced at any stage in the model development pipeline



What is bias? II



Algorithmic bias:

- inadvertent privacy violations
- programmers assign priorities, or hierarchies, for how a program assesses and sorts that data
- collect their own data based on human-selected criteria, which can reflect bias of human designers
- reinforce stereotypes and preferences as they process and display "relevant" data for human users, for example, by selecting information based on previous choices of a similar user or group of users

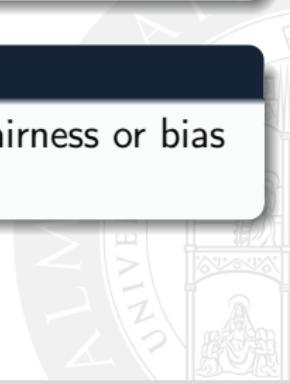
Computational Fairness

Computational fairness

- potential biases and discrimination that can arise from the use of computational algorithms
- ensuring algorithms do not perpetuate or amplify existing biases and do not discriminate against certain groups of people based on sensitive attributes

Fairness Metrics

Quantitative measurement used to assess and quantify the fairness or bias of an algorithm's predictions or decisions



Next in Line...

- 1 Why fairness?
- 2 Goal of the talk
- 3 Fairness and bias in AI
- 4 Fairness Awareness
- 5 Enforcing Fairness
- 6 Our advancements in the field
- 7 Conclusions, Challenges and Opportunities



Fairness Awareness I

Two elements required

- Definition of *fairness notions* (context-dependent, social perspective)
- Quantitative mechanism to measure them

Most approaches based on

→ notion of protected or sensitive variables and (un)privileged groups

- **groups** (defined by one or more sensitive variables) that are disproportionately (less) more likely to be positively classified
- **protected variables** define the aspects of data that are socioculturally precarious for the application of ML
 - gender, ethnicity, age, their synonyms, and essentially any other feature of the data that involves or concerns people

Fairness Awareness II

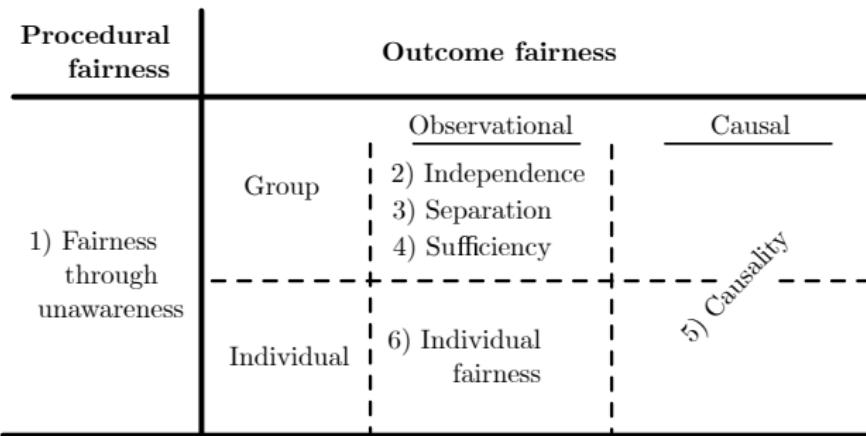


Figure: Organising framework of algorithmic fairness metrics

Procedural Fairness

Procedural Fairness

- concept inherited from administrative law concerned with equality of treatment *within the process* that carries out a decision
- in the computational area
 - not including sensitive attributes in the AI algorithm → omission of sensitive attributes or *fairness through unawareness*
- model accuracy is reduced
- *discrimination effects* do not improve as a consequence of neglecting relationships with proxy
 - ignoring prejudice may not be caused by a single variable but rather by a combination of several ones
- omissions potentially increase bias or discrimination

Procedural fairness	Outcome fairness		
	Group	Observational 2) Independence 3) Separation 4) Sufficiency	Causal 5) Causality 6) Individual fairness
1) Fairness through unawareness	Individual		
			5) Causality

[Bacelar, 2021]

Outcome Fairness

Outcome Fairness: equality of the outcomes (*fair result*)

- two orthogonal groups of two dimensions each:
- *individual* vs. *group* notions of fairness (not mutually exclusive)
- *observational* vs. *causal* approaches

Procedural fairness	Outcome fairness		
	Group	Observational	Causal
1) Fairness through unawareness		1) Independence 2) Separation 3) Sufficiency	
Individual		6) Individual fairness	5) Causality

- *individual* notions of fairness compare single outcomes for individuals
- *group* notions of fairness work on outcomes aggregated by several individuals belonging to the same sensitive category
- *observational*, joint distributions of observable aspects such as outcomes, decisions, features, and sensitive attributes;
- *casual* in case the causal inference is required to acquire knowledge about variables and their (co)relations

Outcome Fairness: Observational Fairness I

Four categories in the set of observational fairness

Procedural fairness	Outcome fairness		
	Group	Observational	Causal
1) Fairness through unawareness		2) Independence 3) Separation 4) Sufficiency	
Individual		6) Individual fairness	5) Causality

- *group* notions of fairness metrics built upon three main abstract fairness criteria
 - 1 independence
 - 2 separation
 - 3 sufficiency
- }
- Consider aspects of a classifier:
- sensitive variable A ,
 - target variable Y and
 - classification score R
- ⇒ a relation of mutual exclusion exists between the three
- 4 *individual* notion of fairness (similarity metric not easy to be defined, computationally infeasible)

Outcome Fairness: Observational Fairness II

Advantages & Limitations

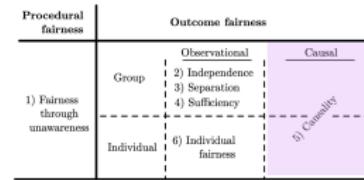
- easiness of state and a lightweight formalism
- assumptions excluded from the inner workings of the classifier, the impact of the decisions, correlations between features and outcomes
- major drawback: limitations in the scope of the evaluation of the available data
- i.e., they do not evaluate what is not observable [Kilbertus et al., 2017]



Outcome Fairness: Causal Fairness

Exploiting the causal graph and the observed data

- enables hidden relationships to be discovered
- identify and mitigate discrimination at its root causes, rather than just on observed disparities



Advantages & Limitations

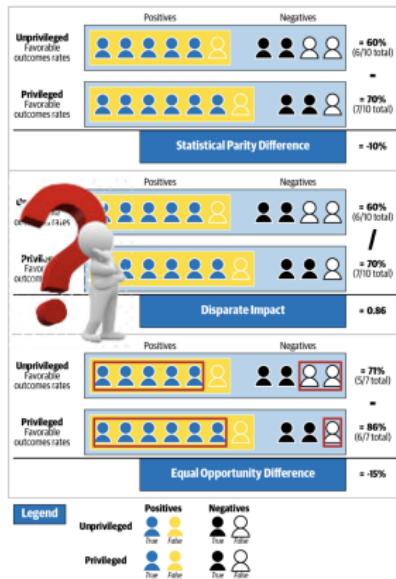
- deeper understanding → more effective interventions
- more precise and effective fairness measures
- fairness by design: building fairness from the design phase
- complexity: deep understanding of causality, domain expertise, access to high-quality data (data availability)
- resource-intensive, in terms of expertise and computational resources
- less interpretable than traditional machine learning models
- sometimes conflicts with legal and regulatory requirements

Which sensitive attributes?

- which variables should be protected?
- which variables are correlated, proxies or quasi-identifiers (when combined identify)?



Which notions of fairness?



- trade-off with accuracy or related metrics
- often conflicting

Looking for properties:

- *incrementally conservative fairness measure*: if the degree to which the measure is satisfied does not decrease if we increase the accuracy of the predictor
- dataset metric
- ...

Next in Line...

- 1 Why fairness?
- 2 Goal of the talk
- 3 Fairness and bias in AI
- 4 Fairness Awareness
- 5 Enforcing Fairness
- 6 Our advancements in the field
- 7 Conclusions, Challenges and Opportunities



Fairness: timing of intervention

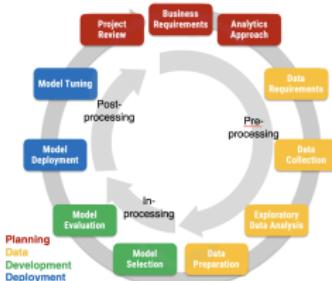


Figure: AI lifecycle & fairness intervention time

- the training data (**pre-processing**):
 - argued to be the most flexible part for repairing bias in the pipeline
 - odds with policies (like GDPR's) potentially introducing new biases
- the learning algorithm (**in-processing**):
 - higher technological effort and integration with ML libraries required
- the predictions (**post-processing**):
 - the accuracy is suboptimal

Fairness intervention techniques in classification

		Procedural	Outcome					
			Group Fairness			Individual Fairness		
			Independence	Separation	Sufficiency	Causality	Causality	Individual fairness
		Fairness through uncertainty	[Feldman et al., 2015]					
	Blinding	[Chen et al., 2020]	[Feldman et al., 2015]					
	Adversarial Learning		[Feng et al., 2014]				[Feng et al., 2014]	
	Causal		[Adu et al., 2018]					
Pre-process	Relabelling					[Gilbertus et al., 2017]	[Kasner et al., 2017]	
						[Mhammed and Chouira, 2021]	[Gupta et al., 2018]	
						[Calders and Verma, 2010]		
	Resampling							
In-process	Resampling					[Avasthi et al., 2021]		
						[Dwork et al., 2018]		
	Resampling							
Adversarial Learning						[Kasner and Calders, 2013]		
						[Calders and Verma, 2010]		
						[Calders and Verma, 2010]		
Post-process	Constraint Optimization					[Ghadri and Starkey, 2019]		
						[Ghadri et al., 2017]		
						[Ghadri et al., 2017]		
						[Ghadri et al., 2018]		
						[Feng et al., 2019]		
	Regularization							
						[Iguelouf et al., 2020]		
						[Zemel et al., 2013]		
						[Corbett-Davies et al., 2017]	[Corbett-Davies et al., 2017]	
						[Alvandi et al., 2021]		
Resampling						[Liu et al., 2019]		
						[Woodworth et al., 2017]		
						[Zafar et al., 2017a]		
Calibration						[Goh et al., 2016]		
						[Quadrato and Shernika, 2017]		
						[Datta et al., 2020]		
						[Alvandi et al., 2021]		
Relabelling						[Datta et al., 2020]		
						[Agapiou et al., 2016]		
Thresholding								

Table: Fairness awareness via fairness notions (columns) and related intervention techniques in the AI lifecycle (rows).

Enforcing Fairness: Challenges

- Algorithm complexity → less interpretable and harder to explain
- Data bias → hard to detect, diversification missing
- Trade-offs balancing fairness and accuracy requires careful consideration
- Resource intensive
- Generalization issues



Next in Line...

- 1 Why fairness?
- 2 Goal of the talk
- 3 Fairness and bias in AI
- 4 Fairness Awareness
- 5 Enforcing Fairness
- 6 Our advancements in the field
- 7 Conclusions, Challenges and Opportunities



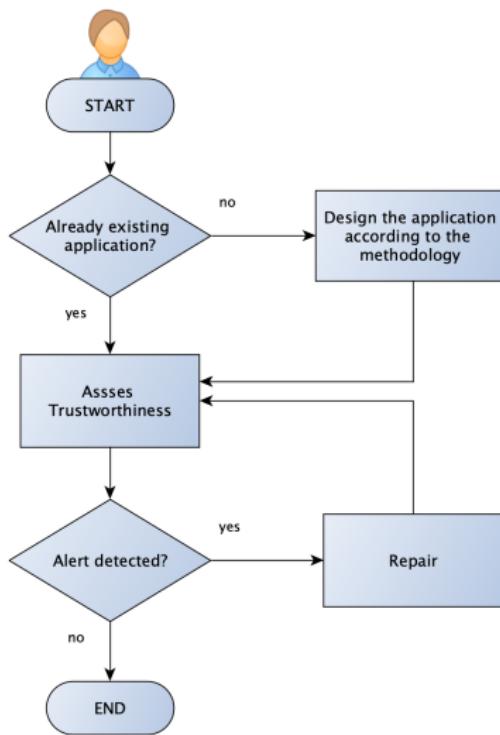
AEQUITAS: Assessment and Engineering of eQuitable, Unbiased, Impartial and Trustworthy Ai Systems

Core idea

*Open controlled experimentation environment for AI stakeholders – provided as a service on the AI on demand platform – to test **fairness** dimensions via **controlled experiments** and to design **trust-by-design** AI applications*

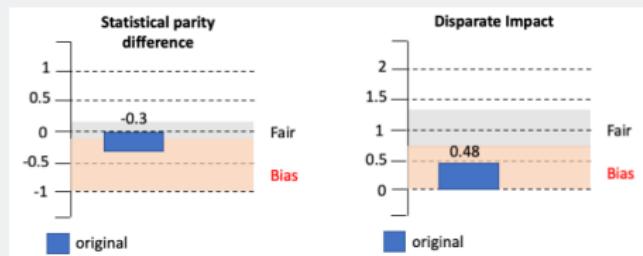


The project idea: workflow



The project idea: an example I

Credit scoring AI application → protected attribute: age (old/young)

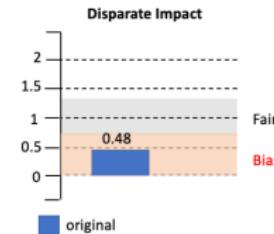
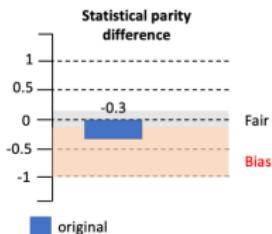


TAI dimension: fairness metrics (*assumptions* to reach fairness)

- **statistical parity difference**: measures the difference that the privileged group get a particular outcome
- **disparate impact**: compares the proportion of individuals that receive a positive output for privileged/unprivileged groups

The project idea: an example II

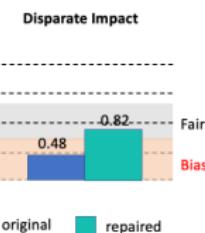
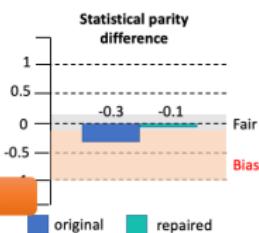
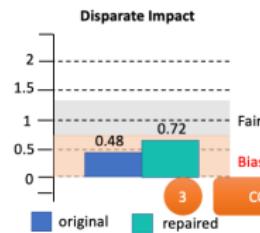
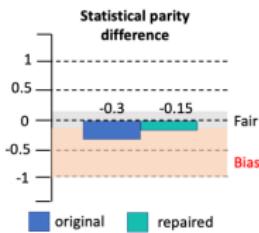
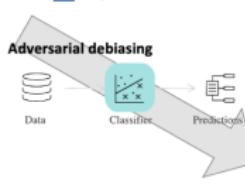
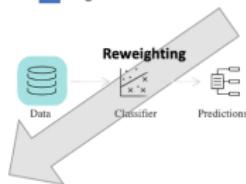
1

ASSESS

2

REPAIR/MITIGATE

- selection of the method
- assumption and conditions



GEOFFair: GEometric Framework for Fairness I

GEometric Framework for Fairness

- represents distributions, ML models, fairness constraints, and hypothesis spaces as vectors and sets
- enables visualization, allowing us to gain insights into the data or the model operation
- enables studying fairness properties in ML



GEOFFair: main definitions

Definition (Ground Vector $y^+ \in \mathcal{Y}^n$)

- data that can be observed and used as ground truth
- paired with the input vector x

Definition (Gold Vector $y^* \in \mathcal{Y}^n$)

- "unbiased" data
- $y^+ = b(y^*)$, where $b : \mathcal{Y}^n \rightarrow \mathcal{Y}^n$ is called the biased mapping

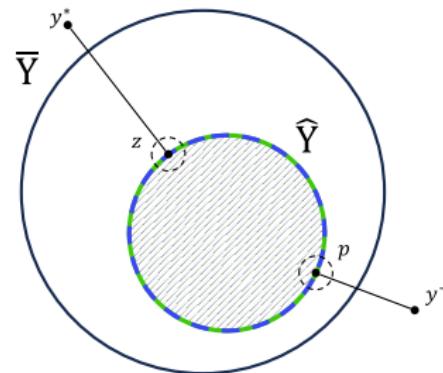
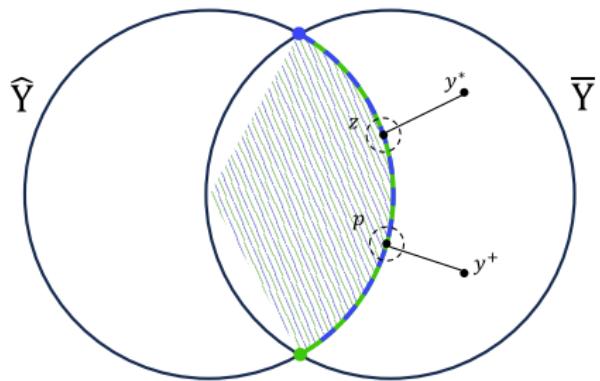
Definition (Hypothesis Space, $\hat{\mathbb{Y}}$)

- set of possible outputs for the chosen class of ML models, i.e.
- $\hat{\mathbb{Y}} = \{y \in \mathcal{Y}^n \mid \exists f \in \mathcal{F} : f(x) = y\}$

Definition (Fair Space, $\bar{\mathbb{Y}} \subseteq \mathcal{Y}^n$)

- set containing all the output vectors aligned with the fairness requirements

GEOFFair: examples



FAiRDAS: Fairness-Aware Ranking as Dynamic Abstract System I

Fairness-Aware Ranking as Dynamic Abstract System

Long-term fairness as an abstract dynamical system

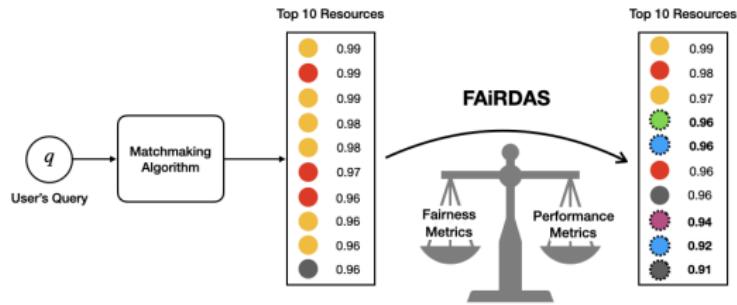
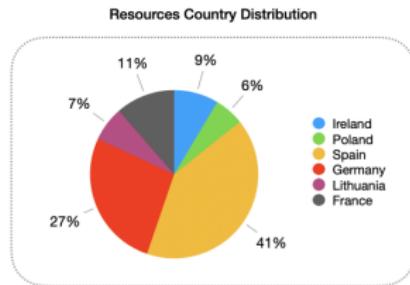
- define metrics of interests (fairness metrics, performance metrics, etc.)
- define the threshold for each metric
- Evolve the system in such a way the metrics remain below the thresholds



FAiRDAS: examples

FAiRDAS: Fairness-Aware Ranking as Dynamic Abstract System

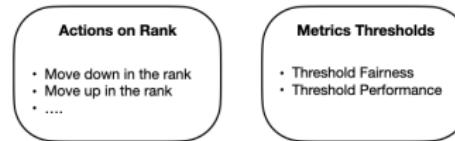
Eleonora Misino, Roberta Calegari, Michele Lombardi, Michela Milano



FAiRDAS Goal



FAiRDAS Key Ingredients



Next in Line...

- 1 Why fairness?
- 2 Goal of the talk
- 3 Fairness and bias in AI
- 4 Fairness Awareness
- 5 Enforcing Fairness
- 6 Our advancements in the field
- 7 Conclusions, Challenges and Opportunities



Conclusions, Challenges and Opportunities I

Which Mechanisms and When?

- legal, ethical, and social context
- selection of the best phase in which to act has dependencies with the data, the availability of the sensitive attributes at testing time, and the fairness notion selected
- context setups can vary between applications

Why Fairness in the AI Lifecycle?

- incorporate fairness needs into the software operations, making it more sustainable from social and technical perspectives
- incorporating fairness seamlessly after the software is operational is in many cases unrealistic given this complexity

Conclusions, Challenges and Opportunities II

Gaps and Challenges

- *Educational aspect* of AI practitioners
- Lack of a *methodological approach* to tackle fairness in the different stages of the AI lifecycle
- *Diversification* is needed beyond existing algorithms and datasets
- *Fairness metrics* need to be balanced between individual and group notions
- *Experimentation environments* are required to provide an easy playground to test different notions and techniques



References |

- [Adel et al., 2019] Adel, T., Valera, I., Ghahramani, Z., and Weller, A. (2019).
One-network adversarial fairness.
In *AAAI*, volume 33, pages 2412–2420.
- [Agarwal et al., 2018] Agarwal, A., Beygelzimer, A., Dudík, M., Langford, J., and Wallach, H. (2018).
A reductions approach to fair classification.
In *International Conference on ML*, pages 60–69. PMLR.
- [Aïvodji et al., 2021] Aïvodji, U., Ferry, J., Gambs, S., Huguet, M.-J., and Siala, M. (2021).
Faircorels, an open-source library for learning fair rule lists.
In *International Conference on Information & Knowledge Management*, pages 4665–4669.
- [Awasthi et al., 2021] Awasthi, P., Beutel, A., Kleindessner, M., Morgenstern, J., and Wang, X. (2021).
Evaluating fairness of machine learning models under uncertain and incomplete information.
In *2021 ACM Conference on Fairness, Accountability, and Transparency*, pages 206–214.
- [Bacelar, 2021] Bacelar, M. (2021).
Monitoring bias and fairness in machine learning models: A review.
ScienceOpen Preprints.

References II

[Bechavod and Ligett, 2017] Bechavod, Y. and Ligett, K. (2017).

Penalizing unfairness in binary classification.

arXiv:1707.00044.

[Beutel et al., 2017] Beutel, A., Chen, J., Zhao, Z., and Chi, E. H. (2017).

Data decisions and theoretical implications when adversarially learning fair representations.

arXiv:1707.00075.

[Calders and Verwer, 2010] Calders, T. and Verwer, S. (2010).

Three naive bayes approaches for discrimination-free classification.

Data mining and knowledge discovery, 21(2):277–292.

[Chen et al., 2019] Chen, J., Kallus, N., Mao, X., Svacha, G., and Udell, M. (2019).

Fairness under unawareness: Assessing disparity when protected class is unobserved.

In *Conference on fairness, accountability, and transparency*, pages 339–348.

[Chiappa, 2019] Chiappa, S. (2019).

Path-specific counterfactual fairness.

In *AAAI*, volume 33, pages 7801–7808.



References III

- [Corbett-Davies et al., 2017] Corbett-Davies, S., Pierson, E., Feller, A., Goel, S., and Huq, A. (2017). **Algorithmic decision making and the cost of fairness.** In *23rd acm sigkdd International Conference on knowledge discovery and data mining*, pages 797–806.
- [Detassis et al., 2020] Detassis, F., Lombardi, M., and Milano, M. (2020). **Teaching the old dog new tricks: supervised learning with constraints.** In *NeHuAI@ ECAI*, pages 44–51.
- [Dwork et al., 2012] Dwork, C., Hardt, M., Pitassi, T., Reingold, O., and Zemel, R. (2012). **Fairness through awareness.** In *3rd innovations in theoretical computer science Conference*, pages 214–226.
- [Dwork et al., 2018] Dwork, C., Immorlica, N., Kalai, A. T., and Leiserson, M. (2018). **Decoupled classifiers for group-fair and efficient machine learning.** In *Conference on fairness, accountability and transparency*, pages 119–133. PMLR.
- [Edwards and Storkey, 2015] Edwards, H. and Storkey, A. (2015). **Censoring representations with an adversary.** *arXiv:1511.05897*.

References IV

[Feldman et al., 2015] Feldman, M., Friedler, S. A., Moeller, J., Scheidegger, C., and Venkatasubramanian, S. (2015).

Certifying and removing disparate impact.

In *21th ACM International Conference on knowledge discovery and data mining*, pages 259–268.

[Feng et al., 2019] Feng, R., Yang, Y., Lyu, Y., Tan, C., Sun, Y., and Wang, C. (2019). Learning fair representations via an adversarial framework.

arXiv:1904.13341.

[Goh et al., 2016] Goh, G., Cotter, A., Gupta, M., and Friedlander, M. P. (2016). Satisfying real-world goals with dataset constraints.

Advances in Neural Information Processing Systems, 29.

[Gupta et al., 2018] Gupta, M., Cotter, A., Fard, M. M., and Wang, S. (2018). Proxy fairness.

arXiv:1806.11212.

[Hardt et al., 2016] Hardt, M., Price, E., and Srebro, N. (2016). Equality of opportunity in supervised learning.

Adv. in neural information processing systems, 29.

References V

- [Ignatiev et al., 2020] Ignatiev, A., Cooper, M. C., Siala, M., Hebrard, E., and Marques-Silva, J. (2020).
Towards formal fairness in machine learning.
In *CP*, pages 846–867. Springer.
- [Kamiran and Calders, 2012] Kamiran, F. and Calders, T. (2012).
Data preprocessing techniques for classification without discrimination.
Knowledge and information systems, 33(1):1–33.
- [Kamiran et al., 2010] Kamiran, F., Calders, T., and Pechenizkiy, M. (2010).
Discrimination aware decision tree learning.
In *2010 IEEE International Conference on Data Mining*, pages 869–874. IEEE.
- [Kamishima et al., 2012] Kamishima, T., Akaho, S., Asoh, H., and Sakuma, J. (2012).
Fairness-aware classifier with prejudice remover regularizer.
In *Joint European Conference on machine learning and knowledge discovery in databases*, pages 35–50. Springer.
- [Kilbertus et al., 2017] Kilbertus, N., Rojas Carulla, M., Parascandolo, G., Hardt, M., Janzing, D., and Schölkopf, B. (2017).
Avoiding discrimination through causal reasoning.
Advances in neural information processing systems, 30.

References VI

- [Krasanakis et al., 2018] Krasanakis, E., Spyromitros-Xioufis, E., Papadopoulos, S., and Kompatsiaris, Y. (2018). Adaptive sensitive reweighting to mitigate bias in fairness-aware classification. In *2018 WWW Conference*, pages 853–862.
- [Kusner et al., 2017] Kusner, M. J., Loftus, J., Russell, C., and Silva, R. (2017). Counterfactual fairness. *Adv. in neural information processing systems*, 30.
- [Leavy et al., 2020] Leavy, S., O’Sullivan, B., and Siapera, E. (2020). Data, power and bias in artificial intelligence. *arXiv:2008.07341*.
- [Leavy et al., 2021] Leavy, S., Siapera, E., and O’Sullivan, B. (2021). Ethical data curation for ai: An approach based on feminist epistemology and critical theories of race. In *AAAI/ACM Conference on AI, Ethics, and Society*, pages 695–703.
- [Liu and Vicente, 2021] Liu, S. and Vicente, L. N. (2021). The sharpe predictor for fairness in machine learning. *arXiv:2108.06415*.

References VII

[Lohia et al., 2019] Lohia, P. K., Ramamurthy, K. N., Bhide, M., Saha, D., Varshney, K. R., and Puri, R. (2019).

Bias mitigation post-processing for individual and group fairness.

In *International Conference on acoustics, speech and signal processing*, pages 2847–2851. IEEE.

[Louizos et al., 2015] Louizos, C., Swersky, K., Li, Y., Welling, M., and Zemel, R. (2015).

The variational fair autoencoder.

arXiv:1511.00830.

[Luong et al., 2011] Luong, B. T., Ruggieri, S., and Turini, F. (2011).

k-nn as an implementation of situation testing for discrimination discovery and prevention.

In *17th International Conference on Knowledge discovery and data mining*, pages 502–510.

[Madras et al., 2018] Madras, D., Creager, E., Pitassi, T., and Zemel, R. (2018).

Learning adversarially fair and transferable representations.

In *International Conference on ML*, pages 3384–3393. PMLR.

[Menon and Williamson, 2018] Menon, A. K. and Williamson, R. C. (2018).

The cost of fairness in binary classification.

In *Conference on Fairness, Accountability and Transparency*, pages 107–118. PMLR.

References VIII

- [Mhasawade and Chunara, 2021] Mhasawade, V. and Chunara, R. (2021). Causal multi-level fairness. In *AAAI/ACM Conf. on AI, Ethics, and Society*, pages 784–794.
- [Pessach and Shmueli, 2021] Pessach, D. and Shmueli, E. (2021). Improving fairness of artificial intelligence algorithms in privileged-group selection bias data settings. *Expert Systems with Applications*, 185:115667.
- [Pleiss et al., 2017] Pleiss, G., Raghavan, M., Wu, F., Kleinberg, J., and Weinberger, K. Q. (2017). On fairness and calibration. *Advances in neural information processing systems*, 30.
- [Quadrianto and Sharmanska, 2017] Quadrianto, N. and Sharmanska, V. (2017). Recycling privileged learning and distribution matching for fairness. *Advances in Neural Information Processing Systems*, 30.
- [Wang et al., 2019] Wang, H., Ustun, B., and Calmon, F. (2019). Repairing without retraining: Avoiding disparate impact with counterfactual distributions. In *International Conference on ML*, pages 6618–6627. PMLR.

References IX

[Woodworth et al., 2017] Woodworth, B., Gunasekar, S., Ohannessian, M. I., and Srebro, N. (2017).

Learning non-discriminatory predictors.

In *Conference on Learning Theory*, pages 1920–1953. PMLR.

[Zafar et al., 2017a] Zafar, M. B., Valera, I., Rodriguez, M. G., and Gummadi, K. P. (2017a). Fairness beyond disparate treatment & disparate impact: Learning classification without disparate mistreatment.

In *26th International Conference on WWW*, pages 1171–1180.

[Zafar et al., 2017b] Zafar, M. B., Valera, I., Rodriguez, M. G., and Gummadi, K. P. (2017b). Fairness constraints: Mechanisms for fair classification.

In *AI and Statistics*, pages 962–970. PMLR.

[Zemel et al., 2013] Zemel, R., Wu, Y., Swersky, K., Pitassi, T., and Dwork, C. (2013). Learning fair representations.

In *International Conference on ML*, pages 325–333. PMLR.